

Multi-Tenancy in Cloud Applications on the Example of PROCEED

by

Felipe Trost

Matriculation Number 456129

A thesis submitted to

Technische Universität Berlin
School IV - Electrical Engineering and Computer Science
Department of Telecommunication Systems
Service-centric Networking

Bachelor's Thesis

March 26, 2025

Supervised by:
Prof. Dr. Axel Küpper

Assistant supervisor:
Kai Grünert

Eidestattliche Erklärung / Statutory Declaration

Hiermit erkläre ich, dass ich die vorliegende Arbeit selbstständig und eigenhändig sowie ohne unerlaubte fremde Hilfe und ausschließlich unter Verwendung der aufgeführten Quellen und Hilfsmittel angefertigt habe.

I hereby declare that I have created this work completely on my own and used no other sources or tools than the ones listed.

Berlin, March 26, 2025

Chuck Norris' son

Abstract

In this thesis, we show that lorem ipsum dolor sit amet.

Zusammenfassung

Hier kommt das deutsche Abstract hin. Wie das geht, kann man wie immer auf Wikipedia nachlesen <http://de.wikipedia.org/wiki/Abstract...>

Contents

1	Introduction	1
1.1	Research Questions	4
1.2	Task List	5
2	Foundations	8
2.1	OAuth 2.0 and OpenID Connect	8
2.1.1	OAuth 2.0 Roles	8
2.1.2	Authorization Grants	9
2.1.2.1	Implicit	9
2.1.2.2	Resource Owner Password Credentials	9
2.1.2.3	Client Credentials	9
2.1.2.4	Authorization Code	9
2.1.3	OpenID Connect	10
2.2	PROCEED's Assets	10
2.3	PROCEED's Role System	11
2.3.1	MS's Role System Terminology	11
2.3.2	MS's Resources and Actions	11
2.3.3	Role Mappings	12
2.3.4	MS's Roles in CASL	12
3	Concept and Design	14
3.1	Modifications to Assets and Resources	14
3.2	Users	14
3.2.1	Authenticated Users and Accounts	15
3.2.2	Merging a Guest User with an Authenticated User	15
3.2.3	Guest User Storage	15
3.2.4	Development Users	16
3.3	Folders	16
3.3.1	Folder Structure Model	16
3.3.2	Storing Assets Inside Folders	19
3.4	Environments	20
3.4.1	Personal Environments	20
3.4.2	Organization Environments	20
3.4.3	Environment Memberships	20

3.4.4	Storing Assets inside Environments	20
3.4.5	Environment Selection	21
3.5	Roles	21
3.5.1	New Resources	22
3.5.2	Enforcing Permissions Based on Folder Structure	23
3.5.3	Default Roles	23
4	Implementation	24
4.1	MS Architecture	24
4.1.1	Endpoints	24
4.1.2	Data storage	25
4.2	Users	25
4.2.1	Sign In Flows	25
4.2.2	Authenticated Users	26
4.2.3	Guest Users	28
4.2.5	Development Users	30
4.3	Assets	30
4.4	Environments	30
4.4.1	Creation of Personal Environments	30
4.4.2	Creation of Organization Environments	31
4.4.3	Adding Users to an Environment	31
4.4.4	Environment Selection	31
4.4.5	Memberships	32
4.5	Roles	32
4.5.1	New Resources	33
4.5.2	Enforcing Roles in the MS	33
4.5.3	Checking Permissions associated to a Folder	34
5	Evaluation	37
6	Conclusion	41
6.1	Outlook	41
	List of Tables	43
	List of Figures	44
	Appendices	47
	Appendix 1	49

1 Introduction

In today's digital age, businesses heavily rely on cloud applications to work on their assets, e.g. documents, spreadsheets, and presentations. Cloud applications are software tools that are accessed and run entirely over the internet. Typically, a copy of the application runs on a remote computer in a data center. Users access it through a web browser or a dedicated application. These tools represent a paradigm shift from traditional software applications, where the majority of the work was done on the user's device. Shifting a part, or even all the workload to a remote computer offers many advantages for both the users and the developers of applications:

- **Accessibility:** Cloud applications can be accessed anywhere from anywhere with an internet connection.
- **Data safety:** There is a big potential for increased data safety, as data can be stored in a secure data center, instead of on a user's device, which could be lost, stolen or damaged.
- **Collaboration:** Collaboration is easier for application developers to implement, as all users could interact with the same instance of the application, instead of having to manage communication between different local applications.
- **Device-agnostic:** Many cloud applications can be accessed from various devices because most of the functionality runs on a server and isn't dependent on the device, unlike traditional software, which often requires a large part of the application to be specifically programmed for each device.
- **Low IT overhead:** Users don't have to set up the application on their own, which would require technical knowledge. And in the case of cloud applications that are accessed through a browser, users don't have to install anything.




Figure 1.1: Users can access cloud applications from different geographical locations and from a variety of devices, such as laptops, smartphones, or tablets, as long as they have an internet connection.

One very common feature that makes these benefits possible is called *multi-tenancy*. *Multi-tenancy* is a software architecture in which an application can be used by multiple users or organizations at the same time. This has to be achieved without a new execution of the application, also referred to as an *instance*, for each user or organization. Without *multi-tenancy*, each user or organization would need to run the application on their own computers, negating most of the benefits listed earlier.

Think of an application that supports *multi-tenancy* like a big apartment building. Each tenant (user or organization) has their own private apartment (their space), where they store their belongings (their assets). Every tenant is living in the same building (cloud application), but each tenant has their own private space.

Many popular cloud applications use this approach. For example, when you use Microsoft Teams, Slack, or Asana, you're sharing the application with many other companies, but you only see and interact with your own team.




Figure 1.2: *Multi-tenancy* in cloud applications: the same instance of the cloud application, can be used by different tenants, with different structures, without them knowing about each other.

The PROCEED Management System (MS for short) is a cloud application for managing business processes. Business processes are sequences of tasks and decision points designed to achieve a specific business goal. PROCEED uses BPMN (Business Process Model and Notation), a standardized graphical notation, to represent these processes visually.

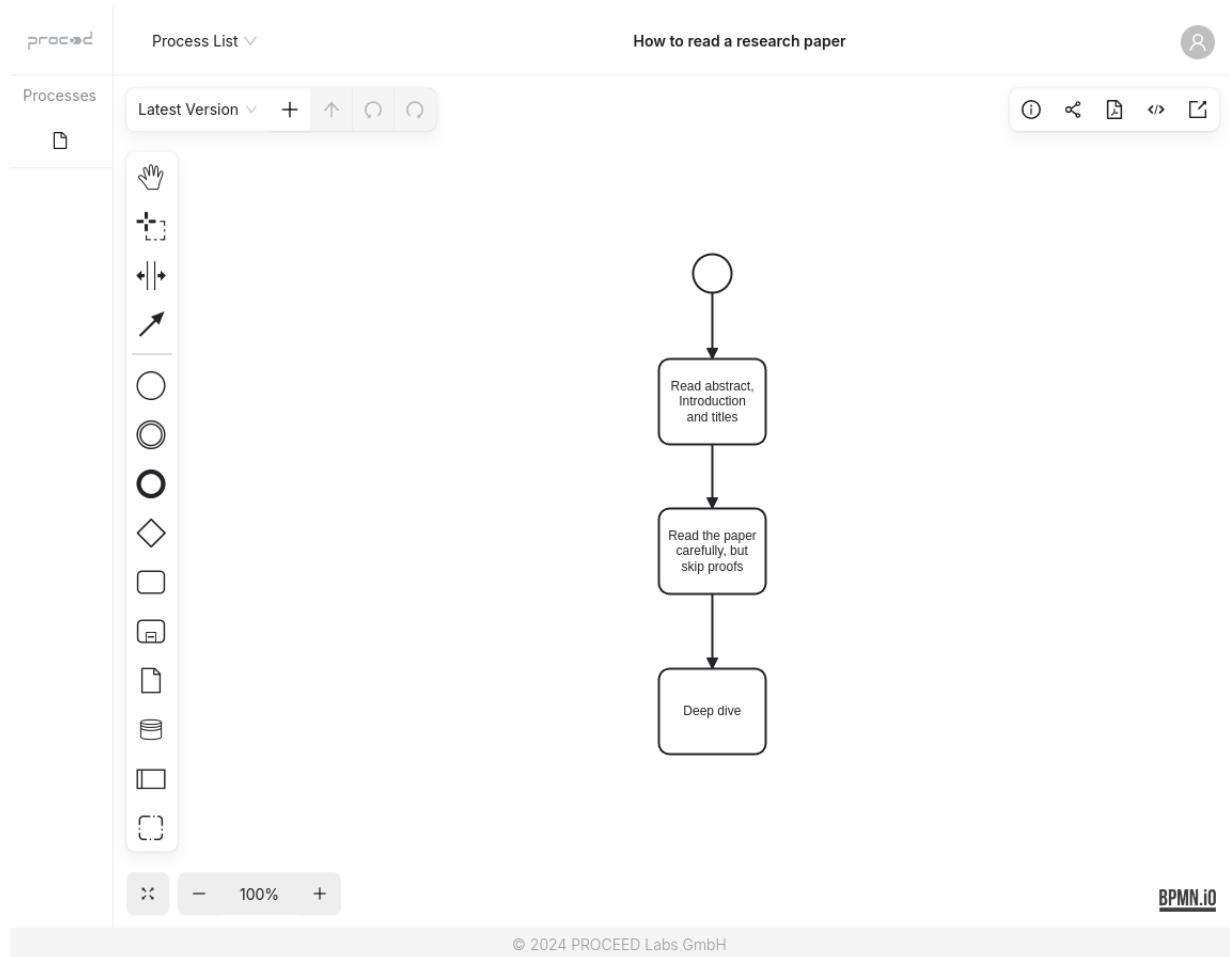


Figure 1.3: Example of process modeling in PROCEED.

Currently, the MS lacks full *multi-tenancy* support. It only supports individual users and doesn't fully support organizations. For organizations to be supported, members of the organization need to be able to have a shared workspace, where they can work on the same assets. However, the MS currently stores all assets in central place, and access to these assets is restricted through a permission system that allows users to see their own assets, or all assets if they possess the admin role. This means that users, regardless of if they are part of an organization or not, can either only see their own assets, or they can see all assets.

For this reason, this thesis will implement *multi-tenant* functionality into the PROCEED MS by introducing the concept of environments. Conceptually, environments are spaces where one or more users can work on shared assets. The MS will be able to hold multiple isolated environments. This way, an organization can have its own environment, where all members of the organization can work on shared assets.




Figure 1.4: Goal of this thesis: tenants with different structures, can work on assets in their own isolated environments in the PROCEED Management System.

1.1 Research Questions

For environments to be successfully integrated into the PROCEED Management System, a few important questions need to be addressed. These questions will help ensure that environments work smoothly with the existing database structure, permission system, and user management. By answering them, we can ensure the implementation is effective

1. Environment representation: How can we integrate environments and their hierarchical folder structure within the MS's storage solution to ensure the following:
 - Data integrity: find a schema that facilitates data consistency after updates.
 - Asset-Environment association: find a schema that associates assets with their respective environments while maintaining a clear separation between multiple environments.
 - Efficiency: find a schema that allows to efficiently query the database.
2. How does PROCEED and its user storage have to be structured to allow users to be members of multiple environments?
3. How can organizations model their hierarchical structures within environments to allow members of the organization to have different levels of access to assets?
4. Users may want to create and manage their own projects independently of an organization: How can users work on personal projects outside an organization?
5. How can environments be implemented in a way that doesn't disrupt existing functionality and data structures?

1.2 Task List

The following task list outlines the concrete steps necessary to address the research questions and ensure a successful implementation of environments. Each task is designed to tackle specific aspects of the MS's architecture, user management, and role assignments.

1. The MS has to support environments, which should be isolated spaces where users can work on assets.
 - a) Every asset in the MS must be stored in only one environment.
 - b) Assets stored in one environment should only be accessed by members of that environment.
 - c) Environments must have a hierarchical folder system to store assets.
 - i. Find a suitable abstraction to represent folders in a database that facilitates consistency after updates and is fast to query.
 - ii. Ensure privacy between environments: folders must only belong to one environment and should only store assets from that environment. Access to folders must be restricted to members of the environment.
2. The MS must be able to hold multiple environments and their assets.
3. The MS must allow different users to access different environments concurrently.
4. Implement personal and organization environments. While both are environments and share common functionality as described in 1, they must behave differently in some situations.
 - a) Personal environments must only have one member.
 - b) Personal environments should only store Processes and Folders.
 - c) Organization environments should have a name and description.
 - d) Organization environments support all assets described in 2.2.
 - e) Organization environments must support multiple members
 - f) Users must be able to create organization environments and invite new members.
 - g) Organization environments must have a role system, where roles can be assigned to users, to manage their access to assets.
 - h) Users of organization environments that have the right permissions must be able to invite users to the organization environment.
5. The MS's user management has to be adapted to fit environments: before the implementation of this thesis, users were strictly tied to one instance of the MS, meaning each time a new instance of the MS was created, a new user storage had to be configured.
 - a) Users have to be global, meaning that they don't belong to any environment.
 - b) Every user has to have a personal environment.
 - c) Personal environments must be tightly coupled with users, i.e. when a user is deleted so is his personal environment.
 - d) Implement Guest users.

- i. The MS must support guest users, which are users that don't sign in with their personal information.
 - ii. The MS should offer a restricted set of functionalities to guest users.
 - iii. Guest users must have the ability to transfer their assets to a normal user.
 - iv. Guest users must not be able to create or be part of organization environments.
6. The MS's preexisting role system must be adapted to fit organization environments and their folder structure: The MS already has a role system in place to manage users' access to resources, this has to be adapted to work with organization environments.
 - a) Roles must belong to only one organization environment.
 - b) The role system must be replicated for each environment, i.e. it works the same as before, with the difference that it is now specific to an environment. E.g. if a role allowed a user to manage all processes before the implementation of environments, now, with the same role, he will be able to modify all processes inside the organization environment the role belongs to.
 - c) Find a suitable permission inheritance model for roles based on the folder structure of an environment (e.g. a user with a role in a parent folder, has the same permissions in all subfolders).
 - d) Roles must always be enforced in the backend, to ensure privacy within organization environments.
 - e) The frontend UI must adapt to a user's roles, by only showing him options that are allowed by his roles.

The following are non-functional tasks that have to be achieved with the implementation of environments in the MS. The goal of these is to ensure that the implementation is user-friendly and most importantly developer-friendly, as many developers will have to work with the codebase in the future. This means that where possible, simple solutions should be favored over complex ones.

1. Keep changes to the MS to a minimum.
2. The user interface for navigating and managing folders and environments should be intuitive and easy to use.
3. Prioritize developer experience by creating clear abstractions and APIs.
 - a) The same data structure and functions should be used for both personal and organization environments where possible. E.g. the same function that creates a folder in a personal environment should be used to create a folder in an organization environment.
 - b) Choose a simple data structure for the folder system, with straightforward functions for modification.
 - c) Streamline environment identification and permissions check in the backend.
 - d) Create simpler helper functions for the frontend of the MS, to adapt the interface to a user's permissions within an environment, than the existing functions for current

role system.

2 Foundations

2.1 OAuth 2.0 and OpenID Connect

OAuth is an open standard for access delegation, commonly used as a way for users to grant client applications access to their information on other applications. OAuth was born as a necessary security measure, to avoid sharing plaintext credentials between applications. Plaintext credential sharing, as outlined in [?] has many security risks:

1. Applications are forced to implement password authentication, to support the sharing of plaintext credentials.
2. Third party applications gain overly broad access to the user's account.
3. Users cannot revoke access to specific third party applications.
4. If any of the third party applications are compromised, the user's account is at risk.

OAuth addresses these issues by decoupling the client application from the role of the resource owner, meaning that the client application will not get a full set of permissions to the user's account. Instead of handing his credentials to the third party application, the resource owner signs in, in the application's website which then issues an access token to the client application. This method avoids the user having to share his credentials with third party applications.

2.1.1 OAuth 2.0 Roles

OAuth 2.0 defines four roles for participants in the protocol flow:

1. Resource owner: The entity that can grant access to a protected resource, typically this would be an end user of a web application.
2. Resource server: The server hosting the protected resources.
3. Client: The application requesting access to the protected resources. OAuth 2.0 distinguishes between two types of clients: confidential and public clients. Confidential clients are capable of keeping their credentials confidential, while public clients, like browser-based applications, cannot.
4. Authorization server: The server that issues access tokens to the client after the resource

owner has been successfully authenticated.

The resource server and the authorization server can be the same entity, but they are not required to be.

2.1.2 Authorization Grants

Authorization Grants are credentials that are issued to clients, which can be exchanged for an access token. This access token can be used to access the protected resources on the resource server. OAuth 2.0 defines four authorization grants with different flows.

2.1.2.1 Implicit

The implicit grant is very helpful for public clients, as it doesn't require confidential client credentials. This is very helpful for browser-based clients, as they can't store confidential credentials securely. In the implicit grant users are redirected to the authorization server, where they authenticate themselves and authorize the client. After which the authorization server issues an access token directly to the client, this is done so with a HTTP redirect, where the access token is embedded in the redirect URL, this way the client can extract the access token from the URL.

In this flow the resource owner only authenticates with the authorization server, thus never having to share his credentials with the client.

Implicit grants have many security risks, as the access token is exposed in the URL and can be intercepted by a malicious attacker. This is why PKCE (Proof Key for Code Exchange) was later introduced as an addition to the implicit grant [?].

2.1.2.2 Resource Owner Password Credentials

This grant type requires the resource owner to share his password credentials with the client. The resource owner's password credentials represent an authorization grant, which the client can exchange them for an access token. Even though this grant type requires the resource owner to share his credentials with the client, these are only used for one request and don't have to be stored.

2.1.2.3 Client Credentials

The client credentials grant is used when the client is the resource owner. Clients are typically issued credentials, which they can use to authenticate themselves. Clients send these credentials to the authorization server and are issued an access token.

2.1.2.4 Authorization Code

The Authorization Code grant is the most common grant type used in OAuth 2.0, it is similar to the implicit grant 2.1.2.1, as it also uses HTTP redirects and it doesn't require the resource

owner to share his credentials with the client. In the authorization code grant, the client redirects the resource owner to the authorization server. There the resource owner authenticates himself and authorizes the client. After which the authorization server redirects the resource owner back to the client with an authorization code. The client then authenticates itself with his confidential credentials on the authorization server and exchanges the authorization code for an access token. As the client needs confidential credentials, this flow is only suitable for confidential clients. The exact steps are shown in figure 2.1.




Figure 2.1: Cloud Application.

2.1.3 OpenID Connect

OpenID (OIDC for short) Connect is an identity layer built on top of the OAuth 2.0. While OAuth 2.0 focuses on authorization (granting clients access to resources), OIDC extends this to authentication. Since OIDC deals with authentication, I will call the resource owner the user from now on.

OIDC uses as its base either the Authorization Code flow 2.1.2.4 or the Implicit flow 2.1.2.1 and it introduces a new type of token called an ID Token. This ID Token is a JSON Web Token (JWT) that contains minimal information about the authenticated user. Most importantly, the ID Token carries a Subject identifier, which is a unique identifier for the user in the <- finish here. The ID Token is sent alongside the Access token to the client, on which step this happens depends on whether the Authorization Code flow or the Implicit flow is used.

In the OIDC flow, after the client obtains the Authorization Code, it exchanges it for both an Access Token (as in OAuth 2.0) and an ID Token. The client can then validate the ID Token to ensure it's genuine and extract the user information contained within.

In essence, OIDC allows clients to obtain information about the authenticated user in a standardized way.

2.2 PROCEED's Assets

Assets are Objects that users can create and manage through the MS's interface. When referring to assets, we are only talking about the core features of the MS, not objects that aid in the usage of the MS, like Roles .e.g., which only helps with managing access to assets. Currently, the MS supports the following assets:

1. Processes, Project and Templates: These objects store BPNN at their core.
2. Task:
3. Machine: Object that represents a server running Distributed Process Engine. The machine is used to manage properties of the server.
4. Execution: An execution represents a process that is being executed distributedly.

Furthermore, the MS implements assets that are used to manage how users can use the MS, called management assets:

- Role2.3: roles are used to manage how users can access assets.
- Rolemapping: role mappings are used to assign roles to users.
- User: represents a user's personal information, e.g. name, username and email.

2.3 PROCEED's Role System

The PROCEED MS uses a Role-Based Access Control system to manage user authorization and determine what actions a user can perform. Roles can be seen as bundles of permissions, which are granted to users. A user can have multiple roles and all the permissions of the roles are additively combined. That is, by adding a permission, a user can never do less than before. Typically, roles are assigned to users based on their job function. RBAC can be advantageous since they can be assigned to multiple users and don't change often, making them easier to manage than individual permissions.

2.3.1 MS's Role System Terminology

The following terms are important to understand the role system in the PROCEED MS:

- Resource: A resource is any protected entity in the management system, that can be accessed by users. Resources can be assets 2.2, but they don't have to be.
- Action: An action is a specific operation that can be performed on a resource.
- Permission: A permission is a tuple of resource type and action, which specifies that a user can perform the action on the resource instances. Optionally a permission can have conditions that have to be met by resource instances, for the user to be able to perform the action.
- Role: A role is a set of permissions. Roles can be assigned to users, which then inherit the role's permissions. Roles can have expiration dates, after which all permissions are revoked.

2.3.2 MS's Resources and Actions

The following are the resource types that are used in the PROCEED MS: Process, Project, Template, Task, Machine, Execution, Role, User, Setting, Rolemapping,

These are the actions that can be performed on these resources: none, view, update, create, delete.

2.3.3 Role Mappings

RoleMappings are a management asset, that is used to assign roles to users. RoleMappings simply store a user ID and a role ID.

2.3.4 MS's Roles in CASL

The PROCEED MS uses ¹CASL to implement Rules. CASL is an isomorphic authorization JavaScript library. To enforce authorization CASL has abilities, which are assigned to users. Abilities expose functions to check whether a user can perform an action on a resource. Abilities are defined by four parameters: user action, subject, fields, conditions. User actions and subjects are analogous to actions and resources 2.3.1.

CASL differentiates between subject type and subject instance. A subject instance is a specific instance of a subject type, e.g. a specific process users are working on, is an instance of the resource type "Process".

Fields are used to specify which fields of a resource instance an action can be performed on, e.g. a user can update a process's name, but not its id, or creation date.

Conditions are used to specify additional conditions that have to be met by a resource instance, for a user to be able to perform an action on it. E.g. a user can only update a process if he created it.

```

1 import { defineAbility } from '@casl/ability';
2
3 class User {
4   constructor(id) {
5     this.id = id;
6   }
7 }
8
9 class Process {
10  constructor(user, name) {
11    this.authorId = user.id;
12    this.createdOn = new Date();
13    this.name = name
14  }
15 }
16
17 function abilityForUser(user) {
18   return defineAbility((can, cannot) => {
19     can('delete', 'User', {id: user.id});
20
21     can('update', 'Process', ['name'], {authorId: user.id});
22   });
23 }
24
25 const user1 = new User(1);
26 const user1Ability = abilityForUser(user1);
27 const user1Process = new Process(user1, 'some process');
28

```

¹ <https://casl.js.org/v6/en/>

```
29 const user2 = new User(2);
30 const user2Ability = abilityForUser(user2);
31
32 user1Ability.can('update', 'Process'); // true
33 user1Ability.can('update', user1Process, 'name'); // true
34 user1Ability.can('update', user1Process, 'createdOn'); // false
35
36 user1Ability.can('delete', user1); // true
37 user1Ability.can('delete', user2); // false
38
39 user2Ability.can('update', 'Process'); // true
40 user2Ability.can('update', user1Process); //false
```

Listing 2.1: CASL example

If there exist any possible resource instance, where the user has permission to perform an action, then the user has permission to perform the action on the resource type. E.g if a user has permission to view some process in the MS, then he has permission to view .

3 Concept and Design

This chapter outlines the key components of the implementation of environments in the MS. The core components are users, environments, roles, and assets. In essence the concept can be summarized as follows: users can be part of multiple environments, which hold Assets. Users that are part of an environment can work on the assets that are stored in it. Each environment has a set of permissions that determine what their users can do with its assets. All other components that will be introduced will help to manage and enforce these relationships.

3.1 Modifications to Assets and Resources

This thesis will modify the assets 2.2 and the resources 2.3.2 supported by the MS. Assets 3.4.4 and Resources 3.5 will be modified to be contained inside environments. Additionally, Environments will be added to the MS' assets and resources. As will be explained in 3.2, Users won't belong to a single environment, they can instead be members of multiple environments, for this reason, users will be removed from the MS's assets and resources. Furthermore, folders 3.3 and memberships 3.4.3 will be added to the MS' assets and resources.

3.2 Users

Previously users in the MS represented a member of a single organization. In order for users to be part of multiple organizations, they can't be tied to a single organization. As a part of the implementation of environments, users are now independent of organizations, they now represent individual people utilizing PROCEED and will be stored as entries in the MS's storage solution 4.1.2.

To facilitate the exploration of the MS, without creating a user, we introduce the option to use the MS as a guest. Thus, we differentiate between authenticated users and guest users. Guest users have a limited feature. Guest users can transition to being an authenticated users whilst retaining their assets, to do this they will need to sign in with personal data.

All users have a personal environment 3.4.1 in which they can create and manage assets freely. Authenticated users can also be part of and create organization environments 3.4.2, where they can collaborate with other users.

3.2.1 Authenticated Users and Accounts

To allow the same user to be able to sign in with different OAuth2 providers, e.g. with Google, Facebook or Discord, we store a separate record, called account, for each of the user's sign-in methods. This means, that the relationship between users and accounts is one-to-many, a user can have multiple accounts, but an account can only be linked to one user. This way, when a user is signing in with credentials from an OAuth2 provider, the MS can look up the corresponding account to the credentials, and then find the user that the account is linked to.




Figure 3.1: Relation between users and accounts.

3.2.2 Merging a Guest User with an Authenticated User

As previously stated, a guest user can transition to being an authenticated user by signing in with his personal data. It could be the case that his personal data already corresponds to an existing user. In this case, the user will be asked if he wants to merge his assets with the existing authenticated user. If he chooses to merge, all the assets in the guest user's personal environment will be transferred to the authenticated user's personal environment. Otherwise, all the assets created by the guest user will be deleted.

3.2.3 Guest User Storage

For storing guest user data, one could take one of two approaches: storing the data in the user's browser or storing it in the MS's database, alongside the data of authenticated users. Storing the data locally has two great benefits: The MS doesn't have to store data of users who might never return and the MS would become less susceptible to an attack where the attacker tries to use up as much space as possible in the MS's storage solution. However, this approach has one key downside, the MS would have to implement two storage solutions and the frontend would need to switch accordingly between them. The added complexity would make it harder for developers to get an overview of the MS's codebase and it also makes it harder to use coding assistance features of IDEs like Goto definition¹. For this reason, storing guest user's data locally isn't worth the benefits. So we decided to store a guest user's data in the MS the same way we do it for authenticated users, with the difference that a flag is set indicate that he is a guest. This way, all the endpoints that authenticated users can call to interact with the MS can also be called by guest users. An important caveat is that, to enforce some of the feature restrictions, relevant endpoints have to check whether the user is a guest. Furthermore, the MS needs to clean up inactive guest users, to prevent the MS's storage solution from filling up with unused data.

¹https://microsoft.github.io/language-server-protocol/specifications/lsp/3.17/specification/#textDocument_definition

3.2.4 Development Users

When developing the MS, it is not common for the developer to have all the necessary keys configured in the environment variables, for the MS to be able to authenticate users with OAuth2 2.1 or send sign-in emails. For this reason the MS implements two development users: johndoe and admin. When the MS is in development mode, i.e. the environment variable `NODE_ENV` is set to `development`, the sign-in page shows a new input field where you can enter the development user's username. Both users are stored as authenticated users in the MS's database and can be used to test the MS during development.

3.3 Folders

Folders will be added to the MS's assets, they are intended to allow organizations to mirror their hierarchical structure and facilitate the organization of assets in general. Starting from a root folder, users should be able to store assets within folders and nest folders inside other folders, creating a flexible and intuitive structure. In this thesis, folders were only implemented to support processes, but they could be extended to support more of the assets that were described in 2.2.

3.3.1 Folder Structure Model


The folder structure is essentially just a rooted tree. The MS doesn't use a database management system, thus it doesn't use a relational language like SQL, still, the data is stored in a way that mimics a relational model. For this reason the rooted folder tree needs to be mapped to a relational model. As outlined by Joe Celko in [?], there are mainly three ways to model trees in relational model: Adjacency List, Nested Set, and Path Enumeration, more commonly known as Materialized Path. Each of these models stores a node as a single entry, but they differ in how they encode their position in the tree.

- **Adjacency List model:** Each node has an identifier and a reference to its parent.
 - Advantages
 - * Finding a node's children is trivial.
 - * Finding a node's parent is trivial.
 - * Adding a node to the tree is trivial and doesn't require updating other nodes.
 - * Moving nodes and their subtrees is trivial, since it only requires updating the parent reference. However, a check has to be made to ensure that the node's new parent isn't one of its descendants, thus creating a cycle.
 - Disadvantages
 - * Finding a node's descendants requires a recursive query, however, this is bounded by the depth of the tree and the amount of nodes.
 - * Finding a node's ancestors, i.e. all the nodes in the path from the root to the node, requires a recursive query, however, this is bounded by the height of the tree and should be efficient.

- * Removing a node requires a recursive query to also delete its descendants, however, this is bounded by the depth of the tree and the amount of nodes.
- **Nested Set model:** Each node has a left and right value, describing an interval starting from the left value and ending at the right value. Children nodes' intervals are contained within the parent's interval. Furthermore, the intervals of siblings are disjoint.
 - Advantages
 - * Finding the descendants of a node is trivial, the query has to select all nodes, whose interval is contained within the node's interval.
 - * Finding a node's ancestors is trivial, the query has to select the all the nodes with a left boundary smaller than the node's left boundary.
 - * Finding a node's parent is trivial, the query has to select the node with the smallest left boundary that is bigger than the node's left boundary.
 - * Deleting a node and its subtree is trivial, the query is analogous to finding the descendants of a node.
 - Disadvantages
 - * Finding a node's children requires a complex query, since it has to select only the node's children without also including the descendants of the children. This query can be inefficient because it has to apply additional filtering to retrieve only the direct children. This adds complexity and can slow down performance.
 - * Adding nodes is non-trivial, since it requires knowledge of its siblings to avoid overlapping intervals. Additionally depending on the implementation, it might require updating intervals of other nodes.
 - * Moving nodes and their subtrees is non-trivial, and inefficient, since it requires updating the intervals of the moved node and all of its descendants. Depending on the implementation it may even be required to update the intervals of other nodes.
- **Materialized Path model:** Each node stores the path from the root to itself.
 - Advantages
 - * Finding a node's ancestors or its parent is trivial, since the path is explicitly stored, the ancestors of a node can be retrieved simply by parsing the stored path.
 - * Finding a node's descendants is easy and efficient, they can be found by querying for nodes whose paths start with the current node's path.
 - * Adding nodes is trivial, as the node's path can be constructed by appending the node's identifier to the parent's path.
 - * Removing a node and its subtree is trivial, as the node's descendants can be easily queried.
 - Disadvantages
 - * Finding a node's children is easy but, not as efficient as finding its descendants, as you need to perform one additional check to ensure that nodes are direct

children.

- * Moving nodes and their subtrees is inefficient, since it requires updating the path of every node in the subtree.
- * Each node stores a string with its path, which can be inefficient in terms of storage space.



Tree

Database models

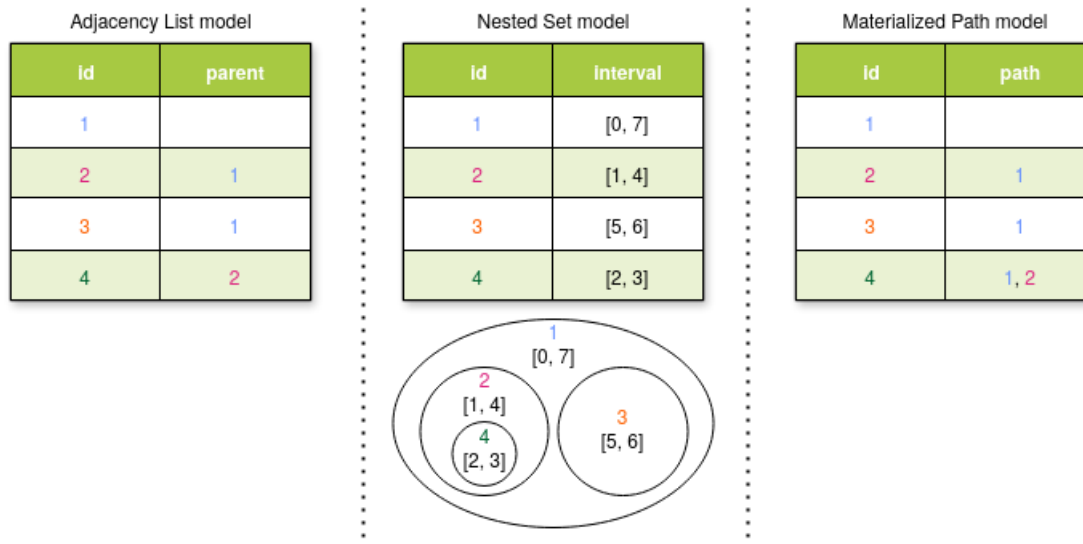


Figure 3.2: Comparison of the Adjacency List, Nested Set, and Materialized Path models.

	Adjacency List	Nested Set	Materialized Path
find children	✓ ✓	X X	✓ ○
find descendants	✓ ○	✓ ✓	✓ ✓
find parent	✓ ✓	X X	✓ ✓
find ancestors	✓ ✓	✓ ✓	✓ ✓
add nodes	✓ ✓	X ○	✓ ✓
remove nodes	✓ ○	✓ ○	✓ ✓
move node	✓ ✓	X X	○ X

✓ _ : Easy query ○ _ : Moderately complex query X _ : Complex query
 _ ✓ : Efficient _ ○ : Moderately efficient _ X : Inefficient

Figure 3.3: Comparison between Adjacency List, Nested Set and Materialized Path models.

To choose the right model, we have to consider the requirements of the MS. Users need to be able to view, add, delete and remove folders. It becomes immediately clear, that the Nested Set model is not suitable for the MS, as adding, removing and moving nodes is inefficient. That leaves us with the adjacency list and materialized path models. The materialized path model has two small advantages: finding descendants and removing a folder together with its subtree is easier. Both queries only require a simple string comparison. However, the adjacency list model isn't far behind on those two points, and is substantially more efficient when it comes to moving nodes. Furthermore, the adjacency list model is better at finding children, which is more valuable to the MS than finding descendants, as the Process view will only show the children. For those reasons, the MS will use the adjacency list model to store the folder structure.

3.3.2 Storing Assets Inside Folders

Once the folder structure is implemented, it still is necessary to store assets inside folders. This thesis only implemented this feature for processes 2.2, however, the same principle could be applied to other assets. A complete redesign of the process' data structures isn't feasible, as it would require rewriting a large part of the MS. For this reason a simple expansion to the data structure was chosen, where analogous to the adjacency list model, each process stores a reference to the folder it is stored in. This way, when moving a folder, it isn't required to update anything other than the folder. Moving an asset to another folder only requires updating the asset's reference.

3.4 Environments

Conceptually, environments are where everything except users are stored. Users aren't stored in environments as they can be a part of multiple environments. Instead, the MS stores memberships, which specify that a user is part of an environment.

There are two types of environments, personal and organization environments. Personal environments are intended for personal use and organization environments are intended for organizations.

3.4.1 Personal Environments

Personal environments are assigned to each user once they sign in. The user for which the environment is created is the only member of this environment, and is therefore called the owner. No other users be a part of this environment. Personal environment only allow users to create and manage processes and folders, while other features offered by the MS can only be used in organization environments 3.4.2.

3.4.2 Organization Environments

Organization environments are intended to be used by organizations, thus they can have a name, description and a logo. In contrast to personal environments, users are allowed to use all the MS' features in an organization environment.

Organization environments can also have multiple Users that are part of it, these are called members.

3.4.3 Environment Memberships

To keep track of which users are part of an environment, a new management asset will be added: memberships. Each membership links one user to one environment, specifying that the user is part of that environment.

3.4.4 Storing Assets inside Environments

All assets within the MS 2.2, including folders, will be modified, so that each asset instance establishes a clear association with a single environment. Every asset will store a reference to the environment it belongs to. This reference is immutable, with the only exception being when assets are transferred from a guest user to an authenticated user 3.2.2.

Processes and folders will be contained within folders, which implies that they belong to the same environment as their root folder. This means that for these assets, we are storing the environment they belong to twice. Storing redundant information is risky as it can lead to inconsistencies if updates aren't done correctly. For instance if a folder's environment reference is changed, but those of its children are not, then the folder structure would span across two environments. This risk is mitigated by the fact, that the environment reference of assets is immutable.

3.4.5 Environment Selection

Users will be a part of multiple environments, and they will need to be able to work on all of them. There are two ways of accomplishing this: the user can select one environment at a time, or he can work on multiple environments at the same time. Environments contain many features and views, making it unfeasible to show them all simultaneously for many environments. Thus, some level of selection is necessary to avoid cluttering the interface.

This selection could be granular, where the user selects per view which environment he wants to work on, however this could lead to confusion, if the user switches views and forgets that he's working on a different environment. For this reason, we chose to have a global environment selection, i.e. all the elements in the UI will show the assets and views of the selected environment.

There are two methods of accomplishing environment selection: implicit and explicit. In the implicit method, the selected environment is managed internally and is not reflected to the user in the URL, this could be accomplished by storing the selected environment in the user's cookies or in the browser's local storage for example. In the explicit method, the environment is encoded in the URL, making the current environment clearly visible. Of course a combination of both methods is also possible, but in order to keep the implementation simple, we only chose one. The implicit method has the advantage that the URLs are shorter and easier to read, however it has the disadvantage, that some links can't be shared, since the implicitly selected environment of another user might not be the same. For this reason, the explicit method was chosen as it allows users to share links with the cost of longer URLs.

3.5 Roles

Roles define what actions a user can perform on an asset. Prior to the changes introduced in this thesis, roles in MS 2.3 were global, meaning their permissions applied universally to all assets across the system. However, with the addition of environments and a folder structure, this approach is no longer practical. Roles will now be tied to a specific organization environment, restricting their permissions to assets within that particular environment. In personal environments, where there is only one user, the user will have full control and be able to perform all actions on his assets without restriction.

Folders allow organizations to mirror their hierarchical structure, but this wouldn't be entirely useful if roles applied to all assets inside the environment. For this reason, roles can now be associated to a folder. A role can define permissions for many assets, of which not all can be stored in folders. If a role is associated with a folder, then, only the permissions that are for assets that can be stored in folders, will be affected by the association. The permissions of roles that are associated with a folder cascade down the folder structure, i.e. a role associated to a folder will also apply to all of its descendants. In this thesis, the folder structure was only implemented for processes, this means that only the permissions for processes and folders will apply to the associated folder's descendants. Roles that aren't associated to a folder will continue to apply to all assets in the environment.




Figure 3.4: Permissions of roles cascade down the folder structure.

If a user's role allows him to view assets and is associated to a folder, then the user also has the permission to view all ancestors of the folder. But this permission is only restricted to the parent folders, not the contents of the parent folders. This allows users to navigate the folder structure until they reach the assets they're allowed to view and manage.

3.5.1 New Resources

With the introduction of environments, all the previously available resources still hold the same meaning, with the difference that they are now associated with an environment. Additionally, two new Resources were introduced, for which roles can specify permissions:

- **Folder:** Folders are used to organize assets in an environment. If a role is associated to a folder, then the permissions for the folder apply to that folder and all folders
 - **view:** Allows the user to view the folder, its name and description. This doesn't necessarily mean that the user can view the assets inside the folder, as he still needs the **view** permission for each children, which would be the case, if the role is associated to a folder.
 - **create:** The user can create new folders inside the folder he has this permission for.
 - **update:** The user can update the folder's name and description, as well as move the folder, to another folder where he has the **create** permission.
 - **delete:** The user can delete the folder.
- **Environment:** Folders are used to organize assets in an environment.
 - **view, create:** These actions are not implemented for environments. Each member can view the name and description of the environment. **create** can't exist as it implies that the environment doesn't exist yet, and therefore a role that contains it cannot exist.
 - **update:** The user can update the environment's name, description and contact number.

- **delete**: Allows the user to delete the environment with all its assets, this is only possible for organization environments,

3.5.2 Enforcing Permissions Based on Folder Structure

In order to enforce permissions based on the folder structure, it is necessary to fetch the newest state of the folder structure every time a user wants to perform an action on an asset. Roles can't store a representation of the folder structure, as it might become outdated. Since permissions need to be verified, for many requests in the MS, and also in the user's browser, to adapt the UI, we decided to compute and cache a representation of the folder structure of every organization environment, from which an asset is requested. This representation is also sent to the user's browser.

3.5.3 Default Roles

For each organization environment two roles will be created, which cannot be deleted and cannot be associated to a folder:

- **@admin**: This role has all permissions for all assets in the organization environment and it is first assigned to the user that creates the organization environment. Only users with the **@admin** role can add new users to this role.
- **@everyone**: The permissions in this role apply to for all the users that are part of the organization environment. The permissions in this role start out empty, but can be modified.

4 Implementation

In this chapter, we will ¹

4.1 MS Architecture

Before diving into the implementation details of environments, it is important to understand the architecture of the MS. The MS is built using Next.js ², a React ³ framework that allows for server-side rendering. Although Next.js' architecture is different from traditional server-side rendered applications and single-page applications, for the purposes of this thesis, it can be thought of as being split into a single-page frontend and a backend. The frontend executes JavaScript code in the user's browser, and is responsible for rendering the UI, handling user input and making requests to the backend. The backend runs on a server and is responsible for handling requests from the frontend, e.g. saving or querying data.

4.1.1 Endpoints

Next.js implements RPC (Remote Procedure Call), by allowing us to write functions that run in the server, which can be imported in the frontend and called with serializable arguments. All files that contain these functions need to use the "use server" directive⁴, either inside the function, or at the top of the file. These functions will be called endpoints from now on.

```
1  // server.js
2  function deleteProcess(environmentId, processId) {
3      "use server"
4
5      // delete process
6  }
7
8  // client.js
9  const button = document.getElementById('delete-process-button');
10 button.addEventListener('click', (e) => {
11     const environmentId = e.target.dataset.environmentId
```

¹<https://github.com/PROCEED-Labs/proceed>

²<https://nextjs.org/>

³<https://reactjs.org/>

⁴<https://react.dev/reference/rsc/use-server>

```
12     const processId e.target.dataset.processId
13     deleteProcess(environmentId, processId);
14   });
```

Listing 4.1: Example of a Zod schema and the corresponding TypeScript type.

4.1.2 Data storage

The MS doesn't use a database management system, instead it stores all data in JSON files. Each file can be seen as a table in a traditional relational database. Even though this approach allows for unstructured data, the MS uses Zod ⁵ schemas to enforce a structure on data that is stored. Zod is a schema declaration and validation library, it allows the MS to define the shape of JSON serializable data. For purposes of simplicity, when we talk about a schema, instead of showing the code that describes the Schema, we will show the typescript type that satisfies the schema.

```
1 import { z } from 'zod';
2
3 const UserSchema = z.object({
4   id: z.string(),
5   username: z.string(),
6   image: z.string().optional(),
7 })
8
9 // TypeScript type that satisfies the UserSchema
10 type User = {
11   id: string;
12   username: string;
13   image?: string | undefined;
14 }
```

Listing 4.2: Example of a Zod schema and the corresponding TypeScript type.

4.2 Users

4.2.1 Sign In Flows

User authentication is implemented by leveraging OpenID Connect2.1.3, with the help of NextAuth.js⁶. A JWT token ⁷ is stored in the user's browser cookies⁸, which is then parsed and verified by the MS' backend. If the JWT token is valid the user is considered authenticated. If the user couldn't be authenticated, he is redirected to the sign-in page.

NextAuth.js has many sign-in methods built-in, which can be set up with little configuration. For email-sign in, NextAuth.js sends an email with a link that the user has to click to sign in. We have to provide a way to store and query the tokens, that are encoded in the link and a

⁵ <https://zod.dev/>

⁶ <https://next-auth.js.org/>

⁷ <https://www.rfc-editor.org/rfc/rfc7519.txt>

⁸ <https://www.rfc-editor.org/rfc/rfc7519.html>

function that sends the email. For OAuth2 providers, we only have to provide a client id and secret.

To store and look up users and accounts, NextAuth.js requires a database adapter, which is a set of functions that interact with the database. The structure of the data that will be stored described in ??.

NextAuth.js provides hooks that allow us to customize the sign-in flow, the most important one for this thesis is the `signIn` hook. This hook is called when a user tries to sign in, this can be a new user or a returning one. As arguments, it receives the user's data, and the account that the user is trying to sign in, if the user doesn't exist the user data may be empty. If the hook returns true, the sign-in flow continues, if it returns a string or false, then the sign-in flow is stopped, and the user is redirected to an error page. Inside this hook, we can also check if the user was previously signed-in as a guest user and is now signing in with personal information, so that we can merge the data of the two, we will elaborate on this in 4.2.4.

```

1 import NextAuth from 'next-auth/next';
2 import GoogleProvider from 'next-auth/providers/google';
3
4 const authOptions = {
5   providers: [
6     EmailProvider({
7       sendVerificationRequest({ url }) {
8         // send email
9       },
10    }),
11    GoogleProvider({
12      clientId: process.env.GOOGLE_CLIENT_ID,
13      clientSecret: process.env.GOOGLE_CLIENT_SECRET,
14    }),
15  ],
16 }
17
18 const handler = NextAuth(authOptions)

```

Listing 4.3: Schema for authenticated users.

4.2.2 Authenticated Users

Authenticated Users are users that sign in to the MS either with their email or with a OAuth 2.0 provider. They're stored in the MS by NextAuth.js after they've successfully signed in. For authenticated users we store an `id`, a flag named `isGuest`, set to false, and personal information. This is the schema for authenticated users:

```

1 {
2   id: string;
3   isGuest: false;
4   emailVerifiedOn: Date | undefined;
5   firstName?: string | undefined;
6   lastName?: string | undefined;
7   username?: string | undefined;
8   image?: string | undefined;

```

```
9     email?: string | undefined;  
10 }
```

Listing 4.4: Schema for authenticated users.

All the personal information is optional, because depending on how the user signs in, the information might not be available. For instance, because NextAuth.js' email sign in only requires the user to input an email, authenticated users are created without a first name, last name or username. A possible workaround would be to automatically generate these, but we chose to leave them undefined, and prompt the user to fill them when he first signs in.




Figure 4.1: Prompt to fill in personal information.

The `image` field is used to store a URL to a user's profile picture. This field can only be set if the user signed in with an OAuth 2.0, which supplied the URL. We don't accept custom URLs, as this allows for attackers to supply URLs to a server he controls, which can be used to track user's browsers. By only saving URLs provided directly by trusted OAuth 2.0 providers, we ensure that the source of the image is reliable.

As stated before, we need Accounts, to store the sign-in methods of a user. To recognize a user's account we need to store the name of the provider and the account's ID on the provider's platform. Additionally, to link the account to a user, we store the user's ID in the account. The schema for accounts is as follows:

```
1 {  
2   id: string;  
3   type: "oauth";  
4   userId: string;  
5   provider: string;  
6   providerAccountId: string;  
7 }
```

Listing 4.5: Schema for accounts.

4.2.3 Guest Users

Users that aren't signed in can choose to try the MS out as a guest, this doesn't require the user to input any personal information. For users that choose to sign in as a guest, a new user is created and stored in the MS. To achieve this, we use a modified version of the authenticated user schema described in 4.4, to only include the `id` and the `isGuest` flag set to `true`. Additionally, we store a reference to an authenticated user, named `signedInWithUserId`, the purpose of this is explained in 4.2.4.

```
1 {  
2   isGuest: true;  
3   id: string;  
4   signedInWithUserId: string | undefined;  
5 }
```

Listing 4.6: Schema for guest users.

This option must be made available to users during the sign-in process. To accomplish this, NextAuth.js provides a built-in `CredentialsProvider`, which enables the implementation of custom sign-in methods. We configured this provider to accept no credentials, allowing users to sign in without supplying personal information.

```
1 import NextAuth from 'next-auth/next';  
2  
3 const authOptions = {  
4   providers: [  
5     ...  
6     CredentialsProvider({  
7       name: 'Continue as Guest',  
8       credentials: {},  
9       async authorize() {  
10        return addUser({ isGuest: true });  
11      },  
12    })),  
13   ],  
14 }  
15  
16 const handler = NextAuth(authOptions)
```

Listing 4.7: Custom sign-in method for guest users.

After a user signs in as a guest, a JWT token is stored in his cookies. Since there is no personal information stored, there is no way for the user to sign in to his guest account from another device. This also means, that if the cookies are deleted, the user will lose access to his guest account.

4.2.4

Guest users can choose to sign in with their email or with an OAuth 2.0 provider. By doing so, their guest account is turned into an authenticated account. The `isGuest` flag is set to `false`, and their sign-in method is stored. All their assets don't need to be transferred, as they are already

stored in the MS, and the user `id` didn't change. This approach only works if the account that the user used to sign in, wasn't already linked to an authenticated user. For the case where the account was already linked to an authenticated user, we store a new value in the guest user's entry, named `signedInWithUserId`, that references the authenticated user that signed in. After signing in, the user is directed to a page, where he can choose to either transfer the guest user's assets to his personal environment, or discard them. This page uses the authenticated user's `id` to see if there are any guest users that signed in with the user's `id`.

```

1 import NextAuth from 'next-auth/next';
2
3 const authOptions = {
4   ...
5   callbacks: {
6     ...
7     signIn: async ({ account, user, email }) => {
8       // Get the user that was signed in when this sign in flow started
9       const session = await getServerSession(nextAuthOptions);
10      const sessionUser = session?.user;
11
12      // Check if the user is signing in
13      if (
14        sessionUser?.guest &&
15        account?.provider !== 'guest-signin' &&
16        !email?.verificationRequest
17      ) {
18        // Check if the user's cookie is correct
19        const sessionUserInDb = getUserById(sessionUser.id);
20        if (!sessionUserInDb || !sessionUserInDb.guest) throw new Error('User id
in session is not valid');
21
22        const userSigningIn = getUserById(user.id);
23
24        if (userSigningIn) {
25          // If the user that is signing in exists, update the guest user
26          updateUser(sessionUser.id, { guest: true, signedInWithUserId:
userSigningIn.id });
27        } else {
28          // If the user that is signing in is a new user, update the guest user
29          updateUser(sessionUser.id, {
30            firstName: user.firstName ?? undefined,
31            lastName: user.lastName ?? undefined,
32            username: user.username ?? undefined,
33            image: user.image ?? undefined,
34            email: user.email ?? undefined,
35            isGuest: false,
36          });
37        }
38      }
39
40      return true;
41    },
42  },
43 }
44
45 const handler = NextAuth(authOptions)

```

Listing 4.8: Handle the transfer of processes from a guest user to an authenticated user.

4.2.5 Development Users

I'm not even sure If this belongs in this thesis

4.3 Assets

- environmentId stored on each thing to improve querying - talk about data normalization - talk about breaking normalization for performance gains -> reference a paper or smth

4.4 Environments

This section will cover the implementation details of environments. In its essence, an environment is just an entry in the MS' storage, where assets point to. To store these environments a new file was added to the MS storage solution 4.1.2 that stores every environment. Every environment has an `id` and a flag named `organization` that indicates what type of environment it is. If an environment is an organization environment, it also stores information about the organization, whereas personal environments store a reference to the user that owns the environment in `ownerId`.

```
1 // Schema for organization environments
2 {
3   id: string;
4   name: string;
5   description: string;
6   organization: true;
7 }
8
9 // Schema for personal environments
10 {
11   id: string;
12   organization: false;
13   ownerId: string;
14 }
```

Listing 4.9: Schemas for organization and personal environments.

4.4.1 Creation of Personal Environments

Personal environments are created when an entry for a user is created, this ensures that every user has a personal environment. Alongside the environment, a new root folder is created for the environment, where the user can store Processes, and other folders. The environment's `ownerId` references the user that created the environment, no more information is necessary, as personal environments are only accessible by one user.

4.4.2 Creation of Organization Environments

Organization environments can be created by signed-in users. In addition to creating the environment itself, two default roles, `@everyone` and `@admin`, are generated. A membership entry is also added to indicate that the creator is part of the environment, and a role mapping is created to assign the creator to the `@admin` role. And a root folder is created for the environment.

Figure 4.2: Form for creating an organization environment.

4.4.3 Adding Users to an Environment

Adding users to an environment is simply done by creating a membership entry, with the organization environment's `id` and the user's `id`. However, the inviter typically doesn't know the invitee's internal `id`, so they can use the invitee's email address instead. The MS then sends an email to the invitee, with an invitation link that contains a token. The link directs to a user to the MS, where the token is verified and the user is added to the organization environment. Additionally, if the inviter has the permissions for it, he can select roles for the new member.

The token in the invitation link is a JWT token, this way there is no need to store it. The token contains the `id`'s of the roles for the invitees, the `id` of the environment, and a reference to the invitee. If the email the inviter provided is associated with a user, then the reference to the invitee stores the user's `id`, otherwise there is the risk, that the invitee changes his email address, and the invitation link is no longer valid for him. If the invitee's email address is not associated with a user, the token just stores his email address. In this case, when the invitee clicks on the link, he is first directed to the sign-in page, and after he signs in he is redirected back to the invitation link.

4.4.4 Environment Selection

The selected environment is encoded in the URL's path like this: `https://staging.proceed-labs.org/<id>/...`. Every view accessed by a user, where URL's path starts with an environment's `id`, will be only related to that environment. Each view can get this environment `id` and fetch the appropriate assets.

```

1 function ProcessesPage(props) {
2   const environmentId = props.params.environmentId;
3
4   const processes = getProcesses(environmentId);
5
6   // render processes
7 }

```

Listing 4.10: Example of a view processes based on the environment id.

Users can switch between environments by selecting them from a dropdown menu, that is placed on the navigation bar.




Figure 4.3: Selection of environments in the MS's navigation bar.

How everything was changed to support environmentst

- memberships - verification (when envs are created by not signed in users) - deletion - managing the env - section for folders - decide how to divide - selection of envs -

Environments are stored as an entry in the MS table

4.4.5 Memberships

4.5 Roles

Each role is now directly associated with an environment by storing an `environmentId` and its permissions only apply to in that environment. Additionally, roles can now be associated to a folder, with the new field `folderId`.

```

1 type Role = {
2   environmentId: string;
3   name: string;
4   permissions: {
5     Process?: number | undefined;
6     Project?: number | undefined;
7     Template?: number | undefined;
8     Task?: number | undefined;
9     Machine?: number | undefined;
10    Execution?: number | undefined;
11    Role?: number | undefined;
12    User?: number | undefined;
13    Setting?: number | undefined;
14    EnvConfig?: number | undefined;
15    RoleMapping?: number | undefined;
16    Share?: number | undefined;
17    Environment?: number | undefined;

```

```

18   Folder?: number | undefined;
19   MachineConfig?: number | undefined;
20   All?: number | undefined;
21 };
22 description?: string | null | undefined;
23 note?: string | null | undefined;
24 expiration?: Date | null | undefined;
25 }

```

Listing 4.11: Example of a view processes based on the environment id.

4.5.1 New Resources

Folders and environments were added to the MS's resource list, and to the schema of Roles, to allow for permissions to be set for them. This way abilities can perform checks on these resources and enforce permissions like described in 2.3.

These resources, were also added to the MS's role management UI:

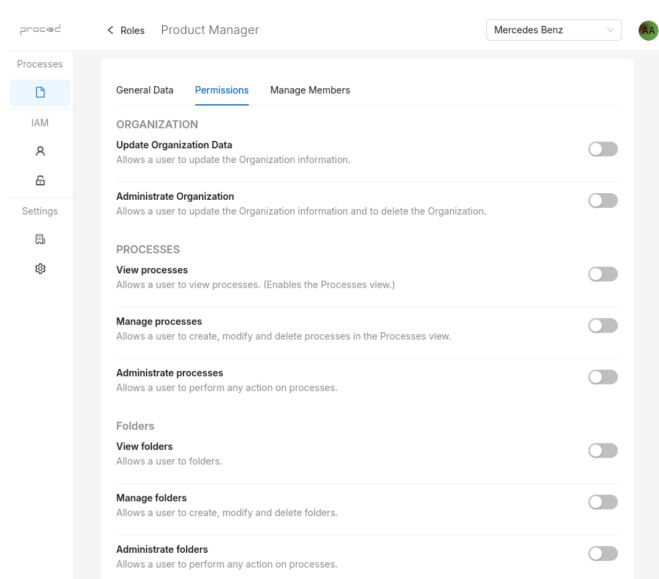


Figure 4.4: Prompt to fill in personal information.

4.5.2 Enforcing Roles in the MS

Every action that a user can perform in an environment, is tied with a permission. For this reason, every endpoint in the MS has to check permissions, fortunately, roles were already implemented in the MS and every endpoint was already using abilities 2.3.4 to check permissions. Abilities are built with the permissions stored in roles that a user has. The only thing that was changed is that every endpoint receives the environment as an argument, and uses the roles the user has in the environment for building his ability.

The function `getCurrentEnvironment` was implemented in order to facilitate getting a user's

ability for an environment, it takes an environment id and uses functionality provided by NextAuth.js to know which user is calling the endpoint.

```

1 async function(processValues, environmentId) {
2   const { ability } = await getCurrentEnvironment(spaceId);
3
4   if (!ability.can('create', toCaslResource('Process', newProcess))) {
5     return userError('Not allowed to create this process');
6   }
7
8   ...
9 }

```

Listing 4.12: Example of an endpoint using an ability.

4.5.3 Checking Permissions associated to a Folder

The process of building abilities was modified to enable them to check a folder structure. Abilities in the MS are built with a conditions matcher ⁹, which is a function that receives a conditions object and returns a function that checks resource instances. The conditions object is defined by each rule, which in turn are derived by the permissions of roles as described in 2.3. If a rule allows an action on a resource instance and includes a conditions object, the rule will only apply if the function returned by the conditions matcher evaluates to **true** when provided with the resource instance.

```

1 import { PureAbility, AbilityBuilder } from '@casl/ability';
2
3 function conditionsMatcher(conditions) {
4   if (conditions.hasToBeAdult) {
5     return (resourceInstance) => resourceInstance.age >= 18;
6   }
7
8   return (resourceInstance) => true;
9 }
10
11 function buildAbility() {
12   const builder = new AbilityBuilder(PureAbility);
13
14   builder.can('view', 'Process', { hasToBeAdult: true });
15
16   return builder.build({ conditionsMatcher });
17 }
18
19 const ability = buildAbility();
20
21 ability.can('read', 'Post', { age: 17 }); // false
22 ability.can('read', 'Post', { age: 23 }); // true

```

Listing 4.13: Simple example of a conditions matcher in CASL.

⁹ <https://casl.js.org/v6/en/advanced/customize-ability#custom-conditions-matcher-implementation>

A simple representation of a folder structure is computed in the MS and stored in a JSON serializable object. Each key is a folder id, and its value is its parent's id. Only the root folder's id isn't contained in the object, as it has no parent. This structure is used by the conditions matcher to check if a process or folder is a descendant or ancestor of a folder. These conditions can be set in the conditions object with the keys `$property_has_to_be_child_of` and `$property_has_to_be_parent_of`. These new conditions are used when turning role's permissions into rules. The conditions matcher also includes a list of seen folders, in the case that the folder structure is has an error and is circular. 4.14 shows a simplified version of the implementation of `$property_has_to_be_child_of`, this implementation doesn't take into account that the conditions object can have multiple conditions.

```

1
2 function conditionsMatcherFactory(folderStructure) {
3   function conditionsMatcher(conditions) {
4     if("$property_has_to_be_child_of" in conditions) {
5       return (resource: any) => {
6         // Folder permissions are also applied to the folder itself
7         if (
8           (resource.__caslSubjectType__ as ResourceType) === 'Folder' &&
9           resource.id === valueInCondition
10        )
11          return true;
12
13        let currentFolder = resource.parentId;
14        const seen = new Set<string>();
15        while (currentFolder) {
16          if (currentFolder === valueInCondition) return true;
17
18          if (seen.has(currentFolder)) throw new Error('Circular reference in
19 folder tree');
20          seen.add(currentFolder);
21
22          // Go up the folder structure
23          currentFolder = folderStructure[currentFolder];
24        }
25
26        return false;
27      };
28    }
29  }
30
31 function buildAbility(folderStructure, rootFolderId) {
32   const builder = new AbilityBuilder(PureAbility);
33
34   builder.can('view', 'Process', {$property_has_to_be_child_of : rootFolderId });
35
36   return builder.build({
37     conditionsMatcher: conditionsMatcherFactory(folderStructure)
38   });
39 }

```

Listing 4.14: Simplified implementation of `$property_has_to_be_child_of` in the conditions matcher.

When a role is being used to build an ability, each permission for each asset is turned into a

rule. When the role is associated to a folder, the rules for processes and folders, use `$property_has_to_be_child_of` in their conditions object, so that the permissions cascade down the folder structure. Additionally, if the role allows a user to view, either a folder or a process, a rule is added, that allows the user to view all ancestors of the asset with `$property_has_to_be_child_of`.

5 Evaluation

In this chapter, we evaluate the implementation of environments in the PROCEED MS by reviewing the requirements defined in section 1.2. The primary goals were to introduce isolated workspaces for both personal and organizational use, enable efficient asset management through a hierarchical folder structure, and adapt the role-based access control system to support these new workspaces. Table 5.1 provides an overview of the task statuses:

Task	Status
1. Implement environments	Partially implemented: folders were only implemented for processes.
2. Personal and organization Environments	Implemented
3. Adapt the MS' user management system to fit environments	Implemented
4. Adapt the MS' role system to fit environments	Implemented

Figure 5.1: Evaluation of Task List.

According to **task 1** Environments are stored as entries in the MS's storage solution with a unique id. Every asset in the MS explicitly stores the id of the environment it belongs to. Memberships to environments are stored in a separate table, where each entry has a user id and an environment id, later on we will explain how the access to assets is managed. Furthermore, each environment has a root folder, which contains more folders and processes. Apart from root folders, both process and folders store the id of the folder they belong to.

According to **task 2** environments store a flag to determine whether they are a personal or an organization environment. Personal environments are created when a user is created and organization environments can be created by signed-in users. Personal environments have a restricted feature-set, this is enforced by the role system, this will be explained in more detail when we address task 4. Users in organization environments can potentially use all the MS's features, this is determined by the roles they have inside that environment. The user that creates an organization environment is assigned the role of admin inside that environment. Access to assets in organization environments is managed by the role system described in task 4, the admin role has all permissions for all assets in the environment. Organization environments can have multiple members, new members can be invited by other members with the right

permissions.

According to **task 3**, the user management was redesigned to accommodate multi-tenancy, users are now stored as records in the MS's storage solution independently of the environments they belong to. This allows for them to be part of multiple environments. Users also have the option to sign in as a guest, which allows them to try the MS inside a personal environment without having to sign up. Guest users are stored as other users, but with a flag that indicates that they are a guest.

According to **task 4**, the MS's role system was adapted to fit environments and folders. In its essence, roles work as they do before with two key differences: Each role belongs to an environment and is only applied for assets inside that environment, and roles can be scoped to folders, meaning that its permissions apply on all descendants of the folder. The role system was also leveraged to restrict the feature set of personal environments: while personal environments don't have roles, the MS still uses the role system to manage access to assets inside personal environments, it gives the owner of the environment all permissions for the allowed features and restricts the rest.

1.2 also describes a set of non-functional requirements mostly related to developer experience. In contrast to functional requirements, which can be evaluated by determining whether the task was completed or not, non-functional requirements are harder to evaluate, as they are subjective. For this reason we determined that surveying the developers that are working on the MS would be the best way to evaluate these non-functional requirements. The survey contained six questions, each question was meant to evaluate one of the non-functional requirements. Most questions were answered with a scale from 1 to 5, for some questions 1 meant that the task was perfectly achieved and for others it was the opposite. For these questions the task is considered as sufficiently met if the average of the answers tends to the favorable side.

Non-functional requirement 1 was to keep the changes to the MS to a minimum, to evaluate this the survey asked: "On a scale of 1 (very minimal) to 5 (extensive), how significantly do you feel the codebase had to change to support multi-tenant environments?". Figure 5.2 shows a histogram of the answers to this question. The answers lie mostly in the middle, with a slight tendency towards the lower end of the scale, thus we conclude that this task was sufficiently met.




Figure 5.2: Histogram of answers to the question: "On a scale of 1 (very minimal) to 5 (extensive), how significantly do you feel the codebase had to change to support multi-tenant environments?"

Non-functional requirement 2 was to provide an intuitive user interface, to evaluate this the

survey asked: "On a scale of 1 (not intuitive at all) to 5 (extremely intuitive), how would you rate the ease of navigating and managing folders/environments in the new UI?". Figure 5.3 shows a histogram of the answers to this question. The answers lie mostly in the upper end of the scale, thus we conclude that the user interface is sufficiently intuitive.




Figure 5.3: Histogram of answers to the question: "On a scale of 1 (not intuitive at all) to 5 (extremely intuitive), how would you rate the ease of navigating and managing folders/environments in the new UI?"

Non-functional requirement 3.a was that personal and organization environments share the same data structures and functions, to improve the developer experience. To evaluate this the survey asked a yes or no question: "Do you feel that personal and organization environments share sufficiently unified structures and functions to use them?". The six developers that answered this question answered "yes", thus we conclude that this task requirement was met.

Non-functional requirement 3.b was to choose a simple data structure for the folder system and to provide simple functions to manage it. To evaluate this the survey asked: "On a scale of 1 (very simple) to 5 (overly complex), how would you rate the complexity of the folder system's data model and its related functions?" Figure 5.4 shows a histogram of the answers to this question. Four people answered "2" and two people answered "3", thus we conclude the user interface is sufficiently intuitive.




Figure 5.4: Histogram of answers to the question: "On a scale of 1 (very simple) to 5 (overly complex), how would you rate the complexity of the folder system's data model and its related functions?"

Non-functional task 3.c was to streamline environment identification and permissions check in the backend. To evaluate this the survey asked: "On a scale of 1 (not at all effective) to 5 (extremely effective), how effective are the new backend helper functions in streamlining en-

vironment identification and permission checks?" Figure 5.5 shows a histogram of the answers to this question. One person voted "3" and the rest of votes were in the favorable side of the scale, thus we conclude that this task was sufficiently met.




Figure 5.5: Histogram of answers to the question: "On a scale of 1 (not at all effective) to 5 (extremely effective), how effective are the new backend helper functions in streamlining environment identification and permission checks?"

Non-functional task 3.d was to introduce simpler helper functions for the frontend, to adapt the interface to a user's permissions within an environment. To evaluate this the survey asked: "On a scale of 1 (no change in complexity) to 5 (drastically simpler), how much do the new frontend helper functions simplify adapting UI elements based on each user's permissions?" Figure 5.6 shows a histogram of the answers to this question. One person voted "3" and the rest of votes were in the favorable side of the scale, thus we conclude that this task was sufficiently met.




Figure 5.6: Histogram of answers to the question: "On a scale of 1 (no change in complexity) to 5 (drastically simpler), how much do the new frontend helper functions simplify adapting UI elements based on each user's permissions?"

6 Conclusion

Building an application structure that supports multi-tenancy presents many challenges, as demonstrated throughout this thesis. In addressing these challenges within the PROCEED Management System (MS), we introduced the concept of environments, isolated workspaces that support both individual and collaborative use cases.

Significant architectural adjustments were required, particularly in restructuring how assets are managed. Assets were adapted to explicitly associate them with specific environments, ensuring clear ownership and secure isolation. To enhance organization and reflect user hierarchies more accurately, a flexible folder structure was implemented. The preexisting role system was adapted to respect environment boundaries and folder structures, ensuring isolation and further enhancing the ability to represent an organization's hierarchy.

Key challenges included restructuring the MS without disrupting existing functionality, ensuring robust privacy and isolation between environments, and adapting the role system to account for both folder and environment-specific permissions. Moreover, the introduction of guest users added complexity to the authentication flows.

6.1 Outlook

While most outlined objectives were successfully achieved, there remain opportunities for future work, such as extending the folder structure to support additional asset types.

Nevertheless, this thesis establishes a robust foundation for multi-tenancy in the PROCEED MS allowing numerous additional features to be explored to further enhance user experience and productivity. For example enabling real-time collaboration on assets such as processes.

Additionally, introducing user directives would significantly simplify user and organization management by enabling automated account creation and streamlined onboarding processes.

List of Tables

List of Figures

1.1	Users can access cloud applications from different geographical locations and from a variety of devices, such as laptops, smartphones, or tablets, as long as they have an internet connection.	2
1.2	<i>Multi-tenancy</i> in cloud applications: the same instance of the cloud application, can be used by different tenants, with different structures, without them knowing about each other.	2
1.3	Example of process modeling in PROCEED.	3
1.4	Goal of this thesis: tenants with different structures, can work on assets in their own isolated environments in the PROCEED Management System.	4
2.1	Cloud Application.	10
3.1	Relation between users and accounts.	15
3.2	Comparison of the Adjacency List, Nested Set, and Materialized Path models. . .	18
3.3	Comparison between Adjacency List, Nested Set and Materialized Path models.	19
3.4	Permissions of roles cascade down the folder structure.	22
4.1	Prompt to fill in personal information.	27
4.2	Form for creating an organization environment.	31
4.3	Selection of environments in the MS's navigation bar.	32
4.4	Prompt to fill in personal information.	33
5.1	Evaluation of Task List.	37
5.2	Histogram of answers to the question: "On a scale of 1 (very minimal) to 5 (extensive), how significantly do you feel the codebase had to change to support multi-tenant environments?"	38
5.3	Histogram of answers to the question: "On a scale of 1 (not intuitive at all) to 5 (extremely intuitive), how would you rate the ease of navigating and managing folders/environments in the new UI?"	39
5.4	Histogram of answers to the question: "On a scale of 1 (very simple) to 5 (overly complex), how would you rate the complexity of the folder system's data model and its related functions?"	39

-
- 5.5 Histogram of answers to the question: "On a scale of 1 (not at all effective) to 5 (extremely effective), how effective are the new backend helper functions in streamlining environment identification and permission checks?" 40
- 5.6 Histogram of answers to the question: "On a scale of 1 (no change in complexity) to 5 (drastically simpler), how much do the new frontend helper functions simplify adapting UI elements based on each user's permissions?" 40

Appendices

Appendix 1

```
1 for($i=1; $i<123; $i++)  
2 {  
3     echo "work harder! ;)";  
4 }
```